CS711008Z Algorithm Design and Analysis
Lecture 10. Algorithm design technique: Network flow and its applications [1]

Dongbo Bu

Institute of Computing Technology
Chinese Academy of Sciences, Beijing, China

Outline

- MAXIMUMFLOW problem: FORD-FULKERSON algorithm, MAXFLOW-MINCUT theorem;
- A duality explanation of FORD-FULKERSON algorithm and MAXFLOW-MINCUT theorem;
- Scaling technique to improve FORD-FULKERSON algorithm;
- Solving the dual problem: Push-Relabel algorithm;
- Extensions of MAXIMUMFLOW problem: lower bound of capacity, multiple sources & multiple sinks, indirect graph;
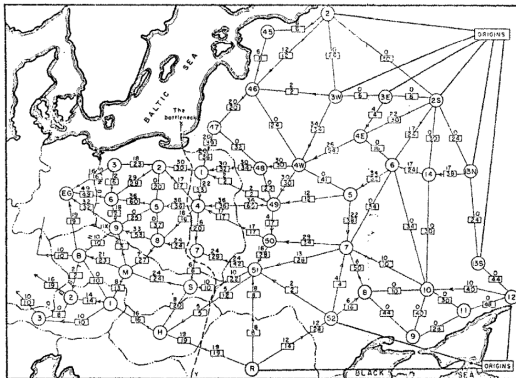
Figure: Soviet Railway network, 1955

- *"From Harris and Ross [1955]: Schematic diagram of the railway network of the Western Soviet Union and Eastern European countries, with a maximum flow of value 163,000 tons from Russia to Eastern Europe, and a cut of capacity 163,000 tons indicated as 'The bottleneck' ..."*

- A recently declassified U.S. Air Force report indicates that the original motivation of minimum-cut problem and Ford-Fulkerson algorithm is to disrupt rail transportation the Soviet Union [A. Shrijver, 2002].

| Year	Developers	Time-complexity
1956	Ford and Fulkerson	$O(mC)$ and $O(m^2 \log C)$
1970	Dinitz	$O(n^2 m)$
1972	Edmonds and Karp	$O(m^2 n)$
1974	Karzanov	$O(n^3)$
1986	Sleator and Tarjan	$O(nm \log n)$
1988	Goldberg and Tarjan	$O(n^2 m \log(\frac{n^2}{m}))$
2012	Orlin	$O(nm)$

MAXIMUMFLOW problem

INPUT:

A directed graph $G = <V, E>$. Each edge e has a capacity C_e.
Two special points: **source** s and **sink** t;

OUTPUT:

For each edge $e = (u, v)$, to assign a flow $f(u, v)$ such that
$\sum_{u,(s,u)\in E} f(s, u)$ is maximized.

Intuition: to push as many commodity as possible from **source** s
to **sink** t.

Flow

Definition (Flow)

$f : E \rightarrow R^+$ is a $s - t$ **flow** if:

1. (Capacity constraints): $0 \leq f(e) \leq C_e$ for all edge e;
2. (Conservation constraints): for any intermediate vertex $v \in V - \{s, t\}$, $f^{in}(v) = f^{out}(v)$, where $f^{in}(v) = \sum_{e \text{ into } v} f(e)$ and $f^{out}(v) = \sum_{e \text{ out of } v} f(e)$. (Intuition: input = output for any intermediate vertex.)

The **value of flow** f is defined as $V(f) = f^{out}(s)$.

Flow and Cut

> **Definition ($s - t$ cut)**
>
> An $s - t$ **cut** is a partition (A, B) of V such that $s \in A$ and $t \in B$.
> The **capacity of a cut** (A, B) is defined as
> $C(A, B) = \sum_{e \text{ from } A \text{ to } B} C(e)$.


$$C(A, B) = 3$$

Flow value lemma

Lemma

(Flow value lemma) Give a flow f. For **any** $s - t$ *cut* (A, B), *the flow across the cut is a constant* $V(f)$. *Formally,*
$V(f) = f^{out}(A) - f^{in}(A).$

$$V(f) = 2 + 0 = 2$$
$$f^{out}(A) - f^{in}(A) = 2 + 1 - 1 = V(f)$$

Proof.

- We have: $0 = f^{out}(v) - f^{in}(v)$ for any node $v \neq s$ and $v \neq t$.
- Thus, we have:

$$
\begin{aligned}
V(f) &= f^{out}(s) - f^{in}(s) \qquad //\text{Hint: } f^{in}(s) = 0; \\
&= \sum\nolimits_{v \in A} (f^{out}(v) - f^{in}(v)) \\
&= (\sum\nolimits_{e \text{ from A to B}} f(e) + \sum\nolimits_{e \text{ from A to A}} f(e)) \\
&\quad - (\sum\nolimits_{e \text{ from B to A}} f(e) + \sum\nolimits_{e \text{ from A to A}} f(e)) \\
&= f^{out}(A) - f^{in}(A)
\end{aligned}
$$

Ford-Fulkerson algorithm [1956]

Figure: Lester Randolph Ford Jr. and Delbert Ray Fulkerson

Trial 1: Dynamic programming technique

- Dynamic programming doesn't seem to work.
- In fact, there is no algorithm known for MAXIMUM FLOW problem that can really be viewed as belonging to the dynamic programming paradigm.
- We know that the MAXIMUMFLOW problem is in $\mathbf{P}$ since it can be formulated as a linear program (See Lecture 8).
- However, the network structure has its own property to enable a more efficient algorithm, informally called **network simplex**.

Trial 2: IMPROVEMENT strategy

- Let's return to the general IMPROVEMENT strategy:

  IMPROVEMENT(f)

  1: $\mathbf{x} = \mathbf{x_0}$; //starting from an initial solution;
  2: **while** TRUE **do**
  3:    $\mathbf{x} =$ IMPROVE($\mathbf{x}$); //move one step towards optimum;
  4:    **if** STOPPING($\mathbf{x}, f$) **then**
  5:       break;
  6:    **end if**
  7: **end while**
  8: **return** $\mathbf{x}$;

Three key questions of iteration framework

- Three key questions:

  1. How to construct an initial solution?

     - For MAXIMUMFLOW problem, an initial solution can be easily obtained by setting $f(e) = 0$ for any e (called 0-flow).
     - It is easy to verify that both CONSERVATION and CAPACITY constraints hold for the 0-flow.

  2. How to improve a solution?
  3. When shall we stop?

- Let p be a simple $s - t$ path in the network G.
  1: Initialize $f(e) = 0$ for all e.
  2: **while** there is an $s - t$ path in graph G **do**
  3:     **arbitrarily** choose an $s - t$ path p in graph G;
  4:     $f = \text{AUGMENT}(p, f)$;
 5: **end while**
  6: **return** f;

- We define $bottleneck(p, f)$ as the minimum residual capacity of edges in path p.

  $\text{AUGMENT}(p, f)$

  1: Let $b = bottleneck(p, f)$;
  2: **for** each edge $e = (u, v) \in p$ **do**
 3:   increase $f(u, v)$ by b;
  4: **end for**

- We start from 0-flow. In order to increase the value of f, we find a $s - t$ path, say $p = s \to u \to v \to t$, to transmit one unit of commodity.

- However we cannot find a $s - t$ path in G to increase f further (left panel) although the maximum flow value is 2 (right panel).


$V(f) = 1$ $V(f) = 2$

- Key observation:
  - When constructing a flow f, one might commit errors on some edges, i.e. the edges should not be used to transmit commodity. For example, the edge $u \to v$ should not be used.

$V(f) = 1$

  - To improve the flow f, we should work out ways to **correct these errors**, i.e. "undo" the transmission assigned on the edges.

- But how to implement the "undo" functionality?
- **Adding backward edges!**
- Suppose we add a **backward** edge $v \to u$ into the original graph. Then we can correct the transmission via pushing back commodity from v to u.

Flow f Backward edges

Definition (Residual Graph)

Given a directed graph $G = <V, E>$ with a flow f, we define **residual graph** $G_f = <V, E'>$. For any edge $e = (u, v) \in E$, two edges are added into E' as follows:

1. (**Forward edge** (u, v) with leftover capacity):
   If $f(e) < C(e)$, add edge $e = (u, v)$ with capacity
   $C(e) = C(e) - f(e)$.
2. (**Backward edge** (v, u) with undo capacity):
   If $f(e) > 0$, add edge $e' = (v, u)$ with capacity $C(e') = f(e)$.

Flow f

Residual graph G_f

Note: the path contains a backward edge (v, u)

Flow f + An $s - t$ path in G_f = New flow f'

Note:

- By using the backward edge $v \to u$, the initial transmission from u to v is pushed back.
- More specifically, the first commodity transferred through flow f changes its path (from $s \to u \to v \to t$ to $s \to u \to t$), while the second one uses the path $s \to v \to t$.

- Let p be a simple $s - t$ path in residual graph G_f, called **augmentation path**. We define $bottleneck(p, f)$ as the minimum capacity of edges in path p.

  FORD-FULKERSON algorithm:
  1: Initialize $f(e) = 0$ for all e.
  2: **while** there is an $s - t$ path in residual graph G_f **do**
  3:     **arbitrarily** choose an $s - t$ path p in G_f;
  4:     $f =$ AUGMENT(p, f);
  5: **end while**
  6: **return** f;

Correctness and time-complexity analysis

Theorem

The operation $f' = \text{AUGMENT}(p, f)$ generates a new flow f' in G.

Flow f + An $s - t$ path in G_f = New flow f'

Proof.

- Checking **capacity constraints**: Consider two cases of edge $e = (u, v)$ in path p:
    1. (u, v) is a forward edge arising from $(u, v) \in E$:
       $0 \leq f(e) \leq f'(e) = f(e) + bottleneck(p, f) \leq$
       $f(e) + (C(e) - f(e)) \leq C(e)$
    2. (u, v) is a backward edge arising from $(v, u) \in E$:
       $C(e) \geq f(e) \geq f'(e) = f(e) - bottleneck(p, f) \geq$
 $f(e) - f(e) = 0$

- Checking **conservation constraints**:
  On each node v, the change of the amount of flow entering v is the same as the change in the amount of flow exiting v.

$\square$

Theorem

(Monotonically increasing) $V(f') > V(f)$

Flow f $+$ An $s-t$ path in G_f $=$ New flow f'

- Hint: $V(f') = V(f) + bottleneck(p, f) > V(f)$ since $bottleneck(p, f) > 0$.

Theorem

$V(f)$ has an upper bound $C = \sum_{e \text{ out of } s} C(e)$.

(Intuition: the edges out of s are completely saturated with flow.)

Property 4: augmentation step

Theorem

Assume all edges have integer capacities. At every intermediate stage of the Ford-Fulkerson algorithm, both flow value $V(f)$ and residual capacities are integers. Thus, $bottleneck(p, f) \geq 1$, and there is at most C iterations of the `while` *loop.*

- Intuition: Under a reasonable assumption that all capacities are integers, we have $bottleneck(p, f) \geq 1$ at every stage; thus, $V(f') \geq V(f) + 1$.
- Time complexity: $O(mC)$.
    - $O(C)$ iterations
    - At each iteration, it takes $O(m + n)$ time to find an $s - t$ path in G_f using DFS or BFS technique.

Theorem

(Tight upper bound) Given a flow f. For any $s - t$ cut (A, B), we have $V(f) \leq C(A, B)$.

$$V(f) = 2 \leq C(A, B) = 3$$

Proof.

$$
\begin{aligned}
V(f) \ &= \ f^{out}(A) - f^{in}(A) \quad &&\text{(by flow value lemma)} \\
&\leq \ f^{out}(A) \quad &&\text{(by } f^{in}(A) \geq 0) \\
&= \ \sum_{e \,\in\, A \,\to\, B} f(e) \\
&\leq \ \sum_{e \,\in\, A \,\to\, B} C(e) \quad &&\text{(by } f(e) \leq C(e)) \\
&= \ C(A, B)
\end{aligned}
$$

□

Theorem

Ford-Fulkerson *ends up with a maximum flow f and a minimum cut (A, B).*

Flow f

Residual graph G_f

Proof.

- FORD-FULKERSON algorithms ends when there is no $s-t$ path in the residual graph G_f.
- Let A be the set of nodes reachable from s in G_f, and $B = V - A$. (A, B) forms a $s-t$ cut. $(A \neq \phi, B \neq \phi)$.
- Consider two types of edges $e = (u, v) \in E$ across cut (A, B):

  1 $u \in A, v \in B$: we have $f(e) = C(e)$ (Otherwise, A should be extended to include v since (u, v) is in G_f.)
  2 $u \in B, v \in A$: we have $f(e) = 0$ (Otherwise, A should be extended to include u since (v, u) is in G_f.)

- Thus we have

$$
\begin{aligned}
V(f) &= f^{out}(A) - f^{in}(A) \\
&= f^{out}(A) \qquad \text{(by } f^{in}(A) = 0\text{)} \\
&= \sum_{e \, \in \, A \to B} f(e) \\
&= \sum_{e \, \in \, A \to B} C(e) \qquad \text{(by } f(e) = C(e)\text{)} \\
&= C(A, B)
\end{aligned}
$$

Ford-Fulkerson algorithm: bad example 1

The integer restriction is important

- In the analysis of FORD-FULKERSON algorithm, the integer restriction of capacities is important: the bottleneck edge leads to an increase of at least 1.
- The analysis doesn't hold if the capacities can be irrational.
- In fact, the flow might be increased by a smaller and smaller number and the iteration will be endless.
- Worse yet, this endless iteration might not converge to the maximum flow.

(See an example by Uri Zwick)

Flow $f : V(f) = 0$

An $s - t$ path in G_f

Flow $f : V(f) = 1$

An $s - t$ path in G_f

Flow $f : V(f) = 2$

Note:

1. After two iterations, the problem is similar to the original problem except for the capacities on $(s, u), (s, v), (u, t), (v, t)$ decrease by 1.

2. Thus, FORD-FULKERSON algorithm will end after 64+32 iterations. (Why? $bottleneck = 1$ at all stages.)

- FORD-FULKERSON algorithm doesn't specify how to choose an augmentation path, leading to some weaknesses:
    - A path with small bottleneck capacity is chosen as augmentation path;
    - We put flow on too many edges than necessary.

- It should be pointed out that the original max-flow paper lists several heuristics for improvement.

- Various strategies to select augmentation path in G_f:
  1. Fat pipes:
     - To select the augmentation path with **the largest bottleneck capacity**;
     - Scaling technique: an efficient way to find an augmentation path with **large** improvement;
  2. Short pipes:
     - Edmonds-Karp: to find **the shortest $s - t$ path** in *BFS tree*.
     - Dinitz' algorithm: to extend *BFS tree* to **layered network**, find augmentation path in the layered network, and perform **amortized** analysis;
     - Dinic's algorithm: running **DFS** to find a collection of augmentation paths (called **blocking flow**) in the **layered network**.

- It should be pointed out that the complexity of EDMONDS-KARP and DINIC'S algorithms do not depend on edge capacities, and thus avoid the limitation of Ford-Fulkerson algorithm for the networks with irrational edge capacities.

Improvement 1: Scaling technique for speed-up

Scaling technique

- Question: can we choose a **large** augmentation path? The larger $bottleneck(p, f)$, the less iterations.

- An $s - t$ path p in G_f with the **largest** $bottleneck(p, f)$ can be found using binary search, or a slight change of Dijkstra's algorithm in $O(m + n \log n)$ time; however, it is still somewhat inefficient.

- Basic idea: we can relax the **"largest"** requirement to **"sufficiently large"**.

- Specifically, we can set up a lower bound Δ for $bottleneck(P, f)$: **simply removing the "small" edges**, i.e. the edges with capacities less than Δ from $G(f)$. This residual graph is called $G_f(\Delta)$.

- Δ will be scaled as iteration proceeds.

- Scaling FORD-FULKERSON algorithm:
  - 1: Initialize $f(e) = 0$ for all e.
  - 2: Let $\Delta = C$;
  - 3: **while** $\Delta \geq 1$ **do**
  - 4:   **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
  - 5:   choose an $s - t$ path p;
  - 6: $f =$ AUGMENT(p, f);
  - 7:   **end while**
  - 8: $\Delta = \Delta/2$;
  - 9: **end while**
  - 10: **return** f;
- Intuition: flow is augmented in a large step size whenever possible; otherwise, the step size is reduced. Step size is controlled via removing the "small" edges out of residual graph.
- Note: Δ turns to be 1 finally; thus no edge in residual graph will be neglected.

Flow $f : V(f) = 0$ No $s - t$ path in G_f

- Flow: 0 flow;
- Δ: $\Delta = 96$;
- $G_f(\Delta)$: the edges in light blue were removed since capcities are less than 96.
- $s - t$ path: cannot find. Thus Δ is scaled: $\Delta = \Delta/2 = 48$.

Flow $f : V(f) = 0$ An $s - t$ path in G_f

- Flow: 0 flow;
- Δ: $\Delta = 48$;
- $G_f(\Delta)$: the edges in light blue were removed since capcities are less than 48.
- $s - t$ path: a path $s - u - t$ appears. Perform augmentation operation.

Flow $f : V(f) = 64$ No $s - t$ path in G_f

- Flow: 64;
- Δ: $\Delta = 48$;
- $G_f(\Delta)$: the edges in light blue were removed since capcities are less than 48.
- $s - t$ path: no path found. Perform scaling: $\Delta = \Delta/2 = 24$.

Flow $f : V(f) = 64$

An $s - t$ path in G_f

- Flow: 64;
- Δ: $\Delta = 24$;
- $G_f(\Delta)$: the edges in light blue were removed since capcities are less than 24.
- $s - t$ path: find a path: $s - v - t$. Perform augmentation.

Flow $f : V(f) = 96$ No $s - t$ path in G_f

- Flow: 96. Maximum flow obtained.
- Δ: $\Delta = 24$;
- $G_f(\Delta)$: the edges in light blue were removed since capcities are less than 24.
- $s - t$ path: cannot find a $s - t$ path.

> **Theorem**
>
> (Outer `while` loop number) The `while` iteration number is at most $1 + \log_2 C$.

SCALING FORD-FULKERSON algorithm:

1: Initialize $f(e) = 0$ for all e.
2: Let $\Delta = C$;
3: **while** $\Delta \geq 1$ **do**
4:      **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
5:          choose an $s - t$ path p;
6:          $f =$ AUGMENT(p, f);
7:      **end while**
8:      $\Delta = \Delta/2$;
9: **end while**
10: **return** f;

Theorem

(Inner `while` loop number) In a scaling phase, the number of augmentations is at most $2m$.

SCALING FORD-FULKERSON algorithm:

1: Initialize $f(e) = 0$ for all e.
2: Let $\Delta = C$;
3: **while** $\Delta \geq 1$ **do**
4:     **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
5:         choose an $s - t$ path p;
6:         $f =$ AUGMENT(p, f);
7:     **end while**
8:     $\Delta = \Delta/2$;
9: **end while**
10: **return** f;

Proof.

1. Let f be the flow that a Δ-scaling phase ends up with, and f^* be the maximum flow. We have $V(f) \geq V(f^*) - m\Delta$. (Intuition: $V(f)$ is not too bad; the distance to maximum flow is small.)

2. In the subsequent $\frac{\Delta}{2}$-scaling phase, each augmentation will increase $V(f)$ at least $\frac{\Delta}{2}$.

Thus, there are at most $2m$ augmentations in the $\frac{\Delta}{2}$-scaling phase. $\square$

- Time-complexity: $O(m^2 \log_2 C)$.
  - $O(\log_2 C)$ outer `while` loop;
  - $O(m)$ inner loops;
  - Each augmentation takes $O(m)$ time.

But why $V(f) \geq V(f^*) - m\Delta$?

Proof.

- Let A be the set of nodes reachable from s in the residual graph $G_f(\Delta)$, and $B = V - A$. Thus (A, B) forms a cut $(A \neq \phi, B \neq \phi)$.

- Consider two types of edges $e = (u, v) \in E$.

  1. $u \in A, v \in B$: we have $f(e) \geq C(e) - \Delta$ (Otherwise, A should be extended to include v since (u, v) in $G_f(\Delta)$.)
  2. $v \in A, u \in B$: we have $f(e) \leq \Delta$ (Otherwise, A should be extended to include v since (u, v) in $G_f(\Delta)$.)

- Thus we have:

$$
\begin{aligned}
V(f) &= \sum_{e \in A \to B} f(e) - \sum_{e \in B \to A} f(e) \\
&\geq \sum_{e \in A \to B} (C(e) - \Delta) - \sum_{e \in B \to A} \Delta \\
&\geq \sum_{e \in A \to B} C(e) - m\Delta \\
&= C(A, B) - m\Delta \\
&\geq V(f^*) - m\Delta
\end{aligned}
$$

Implementation 2: Edmonds-Karp algorithm using $O(m^2n)$ time

Figure: Jack Edmonds, and Richard Karp

Note: The algorithm was first published by Yefim Dinic in 1970 and independently published by Jack Edmonds and Richard Karp in 1972.

EDMONDS-KARP algorithm:

1: Initialize $f(e) = 0$ for all e.
2: **while** there is a $s - t$ path in G_f **do**
3:     choose **the shortest** $s - t$ path p in G_f using BFS;
4:     $f =$ AUGMENT(p, f);
5: **end while**
6: **return** f;

(a demo)

Theorem

Edmonds-Karp algorithm runs in $O(m^2 n)$ time.

Proof.

- During the execution of Edmonds-Karp algorithm, an edge $e = (u, v)$ serves as **bottleneck** edge at most $\frac{n}{2}$ times
- Thus, the while loop will be executed at most $\frac{n}{2}m$ times since there are m edges in total
- It takes $O(m)$ time to find the shortest path using BFS, and augment flow using the path.

$\square$

Lemma

An edge $e = (u, v)$ serves as a *bottleneck* edge at most $\frac{n}{2}$ times.

Proof.

- For a residual graph G_f, we first category all nodes into levels $L_0, L_1, ...,$ where $L_0 = \{s\}$, and L_i contains all nodes v such that the shortest path from s to v has i edges. We use $L(u)$ to denote the level number of node u.

- Consider the two consecutive occurrences of edge $e = (u, v)$ in G_f as bottleneck, say at step k and step k'''.

- We have $L(v) = L(u) + 1$ at step k, and after flow augmentation, the bottleneck edge $e = (u, v)$ will be reversed in G_f.

- At step k''', $e = (u, v)$ becomes a bottleneck edge again.

- This means that $e' = (v, u)$ should be reversed first before step k''', say at step k''. We have $L''(u) = L''(v) + 1$. Thus $L''(u) = L''(v) + 1 \geq L'(v) + 1 \geq L(u) + 2$.

- For any node, its maximal level is at most n and its level never decreases (why?). Thus the lemma holds. □

Step k :

$L(u)$ $L(v)$ $L(v) = L(u) + 1$

s - - - - -> u -> v - - - - -> t

Step k''' :

$L'''(u)$ $L'''(v)$ $L'''(u) \geq L(u) + 2$

s - - - - -> u -> v - - - - -> t

Step k :

$L(u)$ $L(v)$ $L(v) = L(u) + 1$

s ----→ u → v ----→ t

Step $k + 1$:

$L'(u)$ $L'(v)$ $L'(v) \geq L(v)$

s ----→ u ← v ----→ t

$\vdots$

Step k''' :

$L'''(u)$ $L'''(v)$ $L'''(u) \geq L(u) + 2$

s ----→ u → v ----→ t

$L(u)$ $L(v)$ $L(v) = L(u) + 1$

Step k : s u v t

$L'(u)$ $L'(v)$ $L'(v) \geq L(v)$

Step $k + 1$: s u v t

$L''(u)$ $L''(v)$ $L''(u) = L''(v) + 1$

Step k'' : s u v t

$L'''(u)$ $L'''(v)$ $L'''(u) \geq L(u) + 2$

Step k''' : s u v t

Lemma

For any node v, its shortest-path distance $L_f(v)$ in residual graph G_f never decreases during the execution of Edmonds-Karp algorithm.

Node's level never decreases II

Proof.

- Consider a flow f, and the Edmonds-Karp algorithm selects the shortest-path p from s to t in the residual graph G_f for augmentation, forming a new flow f'.

- Assume for contradiction that there is a node v such that $L_{f'}(v) < L_f(v)$

- Without loss of generality, let v be such a node with the minimum $L_{f'}(v)$ in $G_{f'}$. Let $p' = s \rightsquigarrow u \rightarrow v$ be the shortest-path to v in G'_f. Thus we have:

  - $L_{f'}(v) = L_{f'}(u) + 1$
  - $L_{f'}(u) \geq L_f(u)$ due to the selection of v.

- Let's investigate whether $u \rightarrow v$ appears in G_f. There are two cases:

  - $u \rightarrow v \in G_f$
  - $u \rightarrow v \notin G_f$ but $\in G_{f'}$

  and in both cases, a contradiction occurs. Thus the lemma follows.

- $u \to v \in G_f$:
    - We have $L_f(u) + 1 \geq L_f(v)$.
    - Thus $L_{f'}(v) = L_{f'}(u) + 1 \geq L_f(u) + 1 \geq L_f(v)$. A contradiction with the assumption.

- $u \to v \notin G_f$ but $\in G_{f'}$:
    - This can happen only if $v \to u$ exists in p, one of the shortest paths from s to t in G_f.
    - Thus, we have $L_f(u) = L_f(v) + 1$.
    - Thus another contradiction occurs:
      $L_{f'}(v) = L_{f'}(u) + 1 \geq L_f(u) + 1 = L_f(v) + 2$.

Implementation 3: Dinitz' algorithm and its variant Dinic's algorithm

Figure: Yefim Dinitz

The original Dinitz' algorithm

Motivations:

- The initial intention was just to accelerate FORD-FULKERSON algorithm by means of a smart data structure;
- Notice that finding an augmentation path takes $O(m)$ time and becomes a bottleneck of FORD-FULKERSON algorithm;
- It is valuable to save **all information** achieved at a BFS search for subsequent iterations;
- Specifically, the **BFS tree** is enriched to **layered network**:
  - BFS tree: includes **only the first edge found to a node v;**
  - Layered network: keeping **all the edges residing on all the shortest $s - v$ path**. Note that layered network has an advantage to record all shortest $s - t$ paths.

Residual graph G_f BFS tree   Layered network G

- Shimon Even and Alon Itai understood the paper by Y. Dinitz and that by A. Karzanov except for the layered network maintenance (removing the "dead-end" nodes). The gaps were spanned by using:
  1. **blocking flow** (first proposed by A. Karzanov) to prove that the levels of layered network increases from phase to phase;
  2. **DFS** to search an augmentation path guided by the layered network.

- Note: when running on bi-partite graph, the Dinic's algorithm turns into the Hopcroft-Karp algorithm.

1: Initialize $f(e) = 0$ for all e.
2: **while** TRUE **do**
3:     Construct **layered network** G_L from **residual graph** G_f using BFS;
4:     **if** $dist(s,t) = \infty$ **then**
5:         break;
6:     **end if**
7:     find a **blocking flow** f' in G_L using **DFS** technique guided by the layered network;
8:     augment flow f by f';
9: **end while**
10: **return** f;

- The execution of the algorithm can be divided into **phases**, each phase consisting of construction of layered network, and finding blocking flow in it.

- Here, a **blocking flow** refers to a flow such that in the corresponding residual graph, there is no $s - t$ path.

- Intuition: after acquiring a layered network using $O(m)$ time, a blocking flow (containing **a collection of** $s - t$ paths) is found for further augmentation. In contrast, Edmonds-Karp algorithm augment **only one** $s - t$ path.

(a demo here)

Analysis

- Total time: $O(mn^2)$
  - #WHILE $= O(n)$. (why? Each phase increases the level of t by at least 1. Similar argument to the analysis of Edmonds-Karp algorithm.)
  - At each step, it takes $O(m)$ time to construct layered network using extended BFS;
  - and it takes $O(mn)$ time to find blocking flow. The reason is:
    1. it takes $O(n)$ time to find a $s - t$ path using DFS in a layered network;
    2. at least one bottleneck edge in the path will be saturated;
    3. thus it needs at most m iterations to find a blocking flow;

Understanding network-flow from the dual point of view

DUAL: set variables for edges (Intuition: x_i denotes $flow$ via edge i)

$$
\begin{array}{lllllll}
\max & & & & & f & \\
s.t. & x_1 & +x_2 & & & -f & = 0 \text{ vertex } s \\
& & & -x_4 & -x_5 & +f & = 0 \text{ vertex } t \\
& -x_1 & & +x_3 & +x_4 & & = 0 \text{ vertex } u \\
& & -x_2 & -x_3 & & +x_5 & = 0 \text{ vertex } v \\
& x_1 & & & & & \leq C_1 \\
& & x_2 & & & & \leq C_2 \\
& & & x_3 & & & \leq C_3 \\
& & & & x_4 & & \leq C_4 \\
& & & & & x_5 & \leq C_5 \\
& x_1, & x_2, & x_3, & x_4, & x_5 & \geq 0
\end{array}
$$

$$
\begin{array}{llllllll}
\max & & & & & & f & \\
s.t. & x_1 & +x_2 & & & & -f & \leq 0 \text{ vertex } s \\
& & & & -x_4 & -x_5 & +f & \leq 0 \text{ vertex } t \\
& -x_1 & & +x_3 & +x_4 & & & \leq 0 \text{ vertex } u \\
& & -x_2 & -x_3 & & +x_5 & & \leq 0 \text{ vertex } v \\
& x_1 & & & & & & \leq C_1 \\
& & x_2 & & & & & \leq C_2 \\
& & & x_3 & & & & \leq C_3 \\
& & & & x_4 & & & \leq C_4 \\
& & & & & x_5 & & \leq C_5 \\
& x_1, & x_2, & x_3, & x_4, & x_5 & & \geq 0
\end{array}
$$

Note: the constraints (1), (2), (3), and (4) force
$-x_2 - x_3 + x_5 = 0$. So do other constraints.

PRIMAL: set variables for nodes.

$$
\begin{array}{lcccccccc}
\min & & & & C_1z_1 & +C_2z_2 & +C_3z_3 & +C_4z_4 & +C_5z_5 \\
s.t. & y_s & & -y_u & +z_1 & & & & & \geq 0 \\
& y_s & & -y_v & & +z_2 & & & & \geq 0 \\
& & y_u & -y_v & & & +z_3 & & & \geq 0 \\
& -y_t & +y_u & & & & & +z_4 & & \geq 0 \\
& -y_t & & +y_v & & & & & +z_5 & \geq 0 \\
& -y_s & +y_t & & & & & & & \geq 1 \\
& y_s, & y_t, & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & \geq 0
\end{array}
$$

Note:

1. Since the constraints involves the difference among y_s, y_u, y_v and y_t, one of them can be fixed without effects. Here, we fix $y_s = 0$. Thus we have $y_t \geq 1$ (by the constraint $-y_s + y_t \geq 1$).

2. Constraint (4) requires $z_4 \geq y_t - y_u$, and the objective is to minimize a function containing C_4z_4, forcing $y_t = 1$.

3. Constraint (1) requires $z_1 \geq y_u$, and the objective is to minimize a function containing C_1z_1, forcing $z_1 = y_u$. So

An equivalent version

PRIMAL: set variables for nodes.

$$
\begin{array}{llllllll}
\min & & & C_1 z_1 & +C_2 z_2 & +C_3 z_3 & +C_4 z_4 & +C_5 z_5 & \\
s.t. & & -y_u & +z_1 & & & & & = 0 \\
 & & -y_v & & +z_2 & & & & = 0 \\
 & y_u & -y_v & & & +z_3 & & & \geq 0 \\
 & y_u & & & & & +z_4 & & \geq 1 \\
 & & y_v & & & & & +z_5 & \geq 1 \\
 & y_s & & & & & & & = 0 \\
 & & y_t & & & & & & = 1 \\
 & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & \geq 0
\end{array}
$$

Note: the coefficient matrix of constraints (3), (4) and (5) is totally uni-modular, implying the optimal solution is an integer solution.

PRIMAL: set variables for nodes.

$$
\begin{array}{llllllll}
\min & & & C_1 z_1 & +C_2 z_2 & +C_3 z_3 & +C_4 z_4 & +C_5 z_5 \\
s.t. & -y_u & & +z_1 & & & & & = 0 \\
& & -y_v & & +z_2 & & & & = 0 \\
& y_u & -y_v & & & +z_3 & & & \geq 0 \\
& y_u & & & & & +z_4 & & \geq 1 \\
& & y_v & & & & & +z_5 & \geq 1 \\
& y_s & & & & & & & = 0 \\
& y_t & & & & & & & = 1 \\
& y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & = 0/1
\end{array}
$$

$$
\begin{array}{llllllll}
\min & & & C_1 z_1 & +C_2 z_2 & +C_3 z_3 & +C_4 z_4 & +C_5 z_5 & \\
s.t. & & -y_u & +z_1 & & & & & = 0 \\
& & -y_v & & +z_2 & & & & = 0 \\
& & y_u & -y_v & & +z_3 & & & \geq 0 \\
& & y_u & & & & +z_4 & & \geq 1 \\
& & & y_v & & & & +z_5 & \geq 1 \\
& y_s & & & & & & & = 0 \\
& & y_t & & & & & & = 1 \\
& & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & = 0/1
\end{array}
$$

Observations:

- Intuition of primal variables: if node i is in A, $y_i = 0$; and $y_i = 1$ otherwise.
- The primal problem is essentially to find a cut. (Note: $z_1 = 1$ iff $y_s = 0$ and $y_u = 1$, i.e., edge (s, u) is a cut edge.)
- By weak duality, we have $f \leq c$. This is exactly the MaximumFlow-MinimumCut theorem.

FORD-FULKERSON algorithm is essentially a primal-dual algorithm

Primal-dual algorithm

- Recall that the generic primal-dual algorithm can be described as follows.

  1: Initialize $\mathbf{x}$ as a dual feasible solution;
  2: **while** TRUE **do**
  3:    Construct DRP corresponding to $\mathbf{x}$;
  4:    Let ω_{opt} be the optimal solution to DRP;
  5:    **if** $\omega_{opt} = 0$ **then**
  6:       **return** $\mathbf{x}$;
  7:    **else**
  8:       Improve $\mathbf{x}$ according to the optimal solution to DRP;
  9:    **end if**
  10: **end while**

- We will show that solving the DRP is equivalent to finding an augmentation path in residual graph.

- DUAL D: set variables for edges;

$$
\begin{array}{lllllllll}
\max & & & & & & & f & \\
s.t. & x_1 & +x_2 & & & & & -f & \leq 0 \text{ vertex } s \\
& & & & -x_4 & -x_5 & +f & & \leq 0 \text{ vertex } t \\
& -x_1 & & +x_3 & +x_4 & & & & \leq 0 \text{ vertex } u \\
& & -x_2 & -x_3 & & +x_5 & & & \leq 0 \text{ vertex } v \\
& x_1 & & & & & & & \leq 64 \\
& & x_2 & & & & & & \leq 32 \\
& & & x_3 & & & & & \leq 1 \\
& & & & x_4 & & & & \leq 64 \\
& & & & & x_5 & & & \leq 32 \\
& x_1, & x_2, & x_3, & x_4, & x_5 & & & \geq 0
\end{array}
$$

- Let's consider a dual feasible solution $\mathbf{x} = (1, 0, 1, 0, 1)$. Recall how to write DRP from D:
  - Replacing the right-hand side C_i with 0;
  - Adding constraints: $x_i \leq 1$, $f \leq 1$;
  - Keep only the tight constraints J. Here we category J into two sets, i.e. $J = J^S \cup J^E$, where the saturated arcs $J^S = \{i | x_i = C_i\}$, and the empty arcs $J^E = \{i | x_i = 0\}$. Here, $J_S = \{3\}$, and $J_E = \{2, 4\}$.

- DRP:

$$
\begin{array}{rccccccll}
\max & & & & & & f & \\
s.t. & x_1 & +x_2 & & & & -f & = 0 & \text{vertex } s \\
& & & -x_4 & -x_5 & +f & & = 0 & \text{vertex } t \\
& -x_1 & & +x_3 & +x_4 & & & = 0 & \text{vertex } u \\
& & -x_2 & -x_3 & & +x_5 & & = 0 & \text{vertex } v \\
& & & x_i & & & & \leq 0 & i \in J^S \\
& & & x_j & & & & \geq 0 & j \in J^E \\
& x_1, & x_2, & x_3, & x_4, & x_5, & f & \leq 1 &
\end{array}
$$

- $\omega_{OPT} = 0$ implies that optimal solution is found.
- $\omega_{OPT} = 1$ implies a $s - t$ path (with unit flow) in residual graph G_f.


DRP

Residual graph G_f

DRP

Residual graph G_f

- Note that DRP corresponds to finding an augmentation path in the residual graph G_f.
  - $x_i \leq 0, i \in J^S$ denotes a backward edge,
  - $x_j \geq 0, j \in J^E$ denotes a forward edge,
  - and for other edges, there is no restriction for x_i.
- Thus FORD-FULKERSON algorithm is essentially a primal_dual algorithm.

Push-relabel algorithm [A. V. Goldberg, R. E. Tarjan, 1986]

Push-relabel algorithm

The push-relabel algorithm is one of the most efficient algorithms to compute a maximum flow. The general algorithm has $O(n^2m)$ time complexity, while the implementation with FIFO vertex selection rule has $O(n^3)$ running time, the highest active vertex selection rule provides $O(n^2\sqrt{m})$ complexity, and the implementation with Sleator's and Tarjan's dynamic tree data structure runs in $O(nm\log(n^2/m))$ time. In most cases it is more efficient than the Edmonds-Karp algorithm, which runs in $O(nm^2)$ time.

Push-relabel algorithm

- Basic idea: the optimal solution f should meet two constraints simultaneously, namely, f is a flow, and there is no $s - t$ path in the residual graph G_f. FORD-FULKERSON algorithms maintains the first constraint while PUSH-RELABEL maintains the second constraint.

  1. Ford-Fulkerson: set variables for edges. Update flow on edges until G_f has no $s - t$ path;
  2. Push-relabel: set variables for nodes. Update a pre-flow f, maintaining the property that G_f has no $s - t$ path, until f is a flow.

Pre-flow: a relaxation of flow

Definition (Pre-flow)

f is a pre-flow if

- (Capacity condition): $f(e) \leq C(e)$;
- (Excess condition): for any node $v \neq s$,
  $E_f(v) = \sum_{e \text{ into } v} f(e) - \sum_{e \text{ out of } v} f(e) \geq 0$;

- Note that a pre-flow f becomes a flow if no intermediate node has excess.

Label of nodes

Definition (Valid label)

Consider a pre-flow f. A **valid labeling** of nodes is:

- $h(s) = n$, and $h(t) = 0$;
- For each edge (u, v) in the residual graph G_f, we have $h(v) \geq h(u) - 1$;

(Intuition: $h(v)$ is height of the node v, and for an edge in G_f, its end cannot be too lower than its head.)

Theorem

There is no $s - t$ path in a residual graph G_f if there exist valid labels.

Proof.

- Suppose there is a $s - t$ path in G_f.
- Notice that $s - t$ path contains at most $n - 1$ edges.
- Since $h(s) = n$ and $h(u) \leq h(v) + 1$, the height of t should be great than 0. A contradiction with $h(t) = 0$.

□

1: Set f as a pre-flow;
2: Set valid labels for nodes;
3: **while** TRUE **do**
4:   **if** no node has excess **then**
5:     **return** f;
6:   **end if**
7:   Select a node v with excess;
8:   **if** v has a neighbor w such that $h(v) > h(w)$ **then**
9:     **Push** some excess from v to w;
10:   **else**
11:   Perform **relabeling** to increase $h(v)$;
12:   **end if**
13: **end while**

Push-relabel algorithm

1: $h(s) = n$; $h(v) = 0$; for any $v \neq s$;
2: $f(e) = C(e)$ for all $e = (s, u)$; $f(e) = 0$; for other edges;
3: **while** there exists a node v with $E_f(v) > 0$ **do**
4:     **if** there exists an edge $(v, w) \in G_f$ s.t. $h(v) > h(w)$ **then**
5:         //Push excess from v to w;
6:         **if** (v, w) is a forward edge **then**
7:             $e = (v, w)$;
8:             $bottleneck = \min\{E_f(v), C(e) - f(e)\}$;
9:             $f(e)+ = bottleneck$;
10:         **else**
11:             $e = (w, v)$;
12:             $bottleneck = \min\{E_f(v), f(e)\}$;
13:             $f(e)- = bottleneck$;
14:         **end if**
15:     **else**
16: $h(v) = h(v) + 1$; //Relabel node v;
17:     **end if**
18: **end while**

A maximum-flow instance

Initial pre-flow f
and residual graph G_f

Initial pre-flow f
and residual graph G_f

Relabel(u)

Initial pre-flow f
and residual graph G_f

Theorem

Push-relabel algo keeps label valid, and thus outputs a maximum flow when ends.

Proof.

(Induction on the number of push and relabel operations.)

- Push operation: the new f is still a pre-flow since the capacity condition still holds.
  $Push(f, v, w)$ may add edge (w, v) into G_f. We have $h(w) < h(v)$. (pre-condition). Thus, the label is valid for the new G_f.
- Relabel operation: The pre-condition implies $h(v) \leq h(w)$ for any $(v, w) \in G_f$. $relabel(f, h, v)$ changes $h(v) = h(v) + 1$. Thus, the new $h(v) \leq h(w) + 1$.

□

Theorem

For any node v, $\#Relabel \le 2n - 1$. Thus, the total label operation number is less than $2n^2$.

Proof.

1. (Connectivity): For a node w with $E_f(w) > 0$, there should be a path from w to s in G_f.
   (Intuition: node w obtain a positive $E_f(w)$ through a node v by $Push(f, v, w)$. This operation also causes edge (w, v) to be added into G_f. Thus, there should be a path from w to s.)

2. (Upper bound of $h(v)$): $h(v) < 2n - 1$ since there is a path from v to s. The length of the path is less than $n - 1$, $h(s) = n$, and $h(v) \le h(w) + 1$ for any edge (v, w) in G_f.

$\square$

Two types of $Push$ operations:

1. Saturated push (s-push): if $Push(f, v, w)$ causes (v, w) removed from G_f.

2. Unsaturated push (uns-push): other pushes.

$\#Push = \#s\text{-}push + \#uns\text{-}push.$

Theorem

$\#s\text{-}push \leq 2nm.$

Proof.

Consider an edge $e = (v, w)$. We will show that during the execution of algo, (v, w) appears in G_f at most $2n$ times.

- (Removing): a saturated $Push(f, v, w)$ removes (v, w) from G_f. We have $h(v) = h(w) + 1$.
- (Adding): Before applying $Push(f, v, w)$ again, (v, w) should be added to G_f first. The only way to add (v, w) to G_f is $Push(f, w, v)$. The pre-condition of $Push(f, w, v)$ requires that $h(w) \geq h(v) + 1$, i.e., $h(w)$ should be increased at least 2 since the previous $Push(f, v, w)$ operation. And we have $h(w) \leq 2n - 1$.

$\square$

Theorem

$\#uns\text{-}push \leq 2n^2m.$

Time-complexity: $\#Push$ II

Proof.

Define a measure $\Phi(f, h) = \sum_{v:E_f(v)>0} h(v)$.

- (Increase and upper bound) $\Phi(f, h) < 4n^2 m$:
    1. Relabel: a relabel operation increase $\Phi(f, h)$ by 1. The total $O(2n^2)$ relabel operations increase $\Phi(f, h)$ at most $O(2n^2)$.
    2. Saturized push: A saturated $Push(f, v, w)$ operation increases $\Phi(f, h)$ by $h(w)$ since w has excess now. $h(w) \leq 2n - 1$ implies an upper bound for each operation. The total $2nm$ saturated pushes increase $\Phi(f, h)$ by at most $4n^2 m$.

- (Decrease) An unsaturated $Push(f, v, w)$ will reduce $\Phi(f, h)$ at least 1.
  (Intuition: after unsaturated $Push(f, v, w)$, we have $E_f(v) = 0$, which reduce $h(v)$ from $\Phi(f, h)$; on the other side, w obtains excess from v, which will increase $\Phi(f, h)$ by $h(w)$. From $h(v) \leq h(w) + 1$, we have that $\Phi(f, h)$ reduces at least 1.)