

# Introduction to CUDA Programming

Lecture 3: control flow and synchronisation

高性能计算机研究中心

# Warp divergence

- Threads are executed in warps of 32, with all threads in the warp executing the same instruction at the same time
- What happens if different threads in a warp need to do different things?

```
if (x < 0.0)
    z = x - 2.0;
else
    z = sqrt(x);
```

- This is called *warp divergence* – CUDA will generate correct code to handle this, but to understand the performance you need to understand what CUDA does with it

# Warp divergence

This is not a new problem.

Old CRAY vector supercomputers had a logical merge vector instruction

```
z = p ? x : y;
```

which stored the relevant element of the input vectors  $x$ ,  $y$  depending on the logical vector  $p$

```
for(i = 0; i < I; i++) {  
    if (p[i])  
        z[i] = x[i];  
    else  
        z[i] = y[i];  
}
```

# Warp divergence

Similarly, NVIDIA GPUs have *predicated* instructions which are carried out only if a logical flag is true.

```
p: a = b + c; // computed only if p is true
```

In the previous example, all threads compute the logical predicate and two predicated instructions

```
p = (x < 0.0);  
p: z = x - 2.0; // single instruction  
!p: z = sqrt(x);
```

# Warp divergence

## ■ Note that:

- `sqrt(x)` would usually produce a NaN when  $x < 0$ , but it's not really executed when  $x < 0$  so there's no problem
- all threads execute both conditional branches, so execution cost is sum of both branches  $\Rightarrow$  potentially large loss of performance

# Warp divergence

## ■ Another example:

```
if (n >= 0)
    z = x[n];
else
    z = 0;
```

- $x[n]$  is only read here if  $n \geq 0$
- Don't have to worry about illegal memory accesses when  $n$  is negative

# Warp divergence

- If the branches are **big**, nvcc compiler inserts code to check if all threads in the warp take the same branch (*warp voting*) and then branches accordingly.

```
p = ...  
  
    if (any(p)) {  
p:    ...  
p:    ...  
    }  
  
    if (any(!p)) {  
!p:   ...  
!p:   ...  
    }
```

# Warp divergence

## Note:

- doesn't matter what is happening with other warps – each warp is treated separately
- if each warp only goes one way that's very efficient
- warp voting costs a few instructions, so for very simple branches the compiler just uses predication without voting



# Warp divergence

In some cases, can determine at compile time that all threads in the warp must go the same way

e.g. if case is a run-time argument

```
if (case == 1)
    z = x * x;
else
    z = x + 2.3;
```

In this case, there's no need to vote

# Warp divergence

- Warp divergence can lead to a big loss of parallel efficiency – one of the first things I look out for in a new application.
- In worst case, effectively lose factor 32x in performance if one thread needs expensive branch, while rest do nothing
- Typical example: PDE application with boundary conditions
  - if boundary conditions are cheap, loop over all nodes and branch as needed for boundary conditions
  - if boundary conditions are expensive, use two kernels: first for interior points, second for boundary points

# Warp divergence

- Another example: processing a long list of elements where, depending on run-time values, a few require very expensive processing
- GPU implementation:
  - first process list to build two sub-lists of “simple” and “expensive” elements
  - then process two sub-lists separately
- Note: none of this is new – this is what we did 20 years ago on CRAY and Thinking Machines systems.
- What's important is to understand hardware behaviour and design your algorithms / implementation accordingly

# Synchronisation

- Already introduced `__syncthreads()`; which forms a barrier – all threads wait until every one has reached this point.
- When writing conditional code, must be careful to make sure that all threads do reach the `__syncthreads()`;
- Otherwise, can end up in *deadlock*

# Typical application

```
// load in data to shared memory
...
...
...

// synchronisation to ensure this has finished
__syncthreads();

// now do computation using shared data
...
...
...
```

# Synchronisation

- **There are other synchronisation instructions which are similar but have extra capabilities:**
  - `int __syncthreads__ count(predicate)`  
**counts how many predicates are true**
  - `int __syncthreads__ and(predicate)`  
**returns non-zero (true) if all predicates are true**
  - `int __syncthreads__ or(predicate)`  
**returns non-zero (true) if any predicate is true**
- **I've not used these, and don't currently see a need for them**

# Warp voting

- There are similar *warp voting* instructions which operate at the level of a warp:
  - `int __all(predicate)`  
returns non-zero (true) if all predicates in warp are true
  - `int __any(predicate)`  
returns non-zero (true) if any predicate is true
  - `unsigned int __ballot(predicate)`  
sets  $n^{th}$  bit based on  $n^{th}$  predicate
- Again, I've never used these

# Atomic operations

- Occasionally, an application needs threads to update a counter in shared memory.

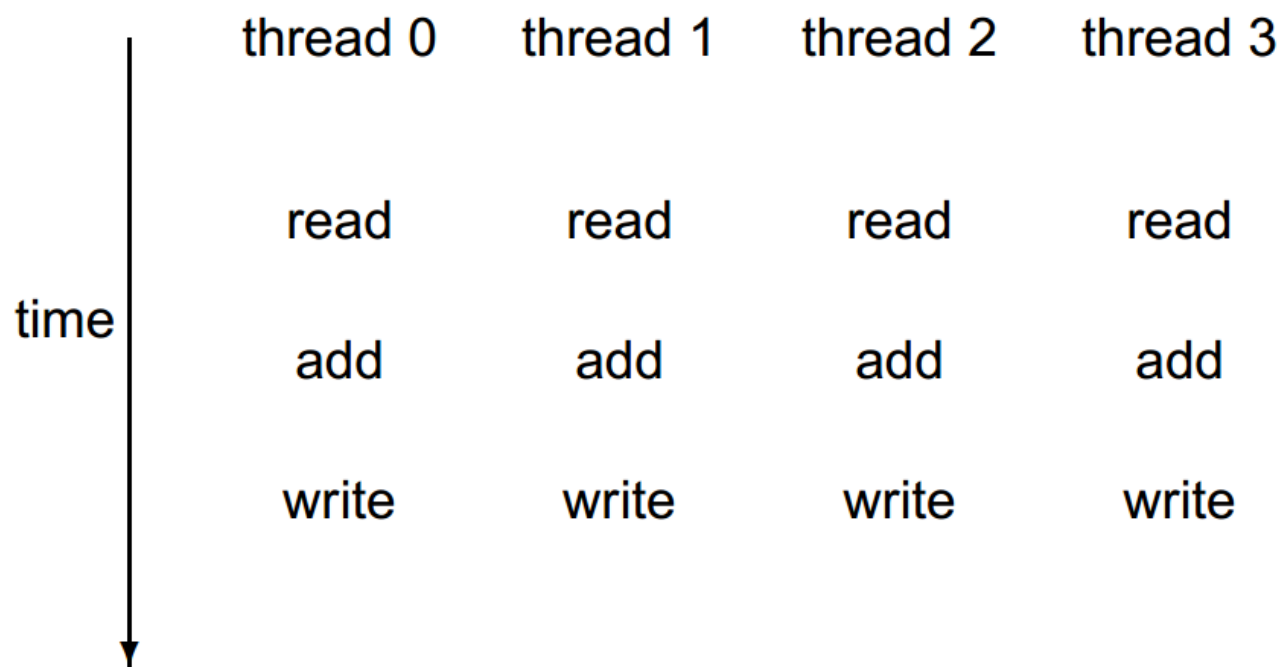
```
__shared__ int count;  
  
...  
  
if ( ... )  
    count++;
```

- In this case, there is a problem if two (or more) threads try to do it at the same time



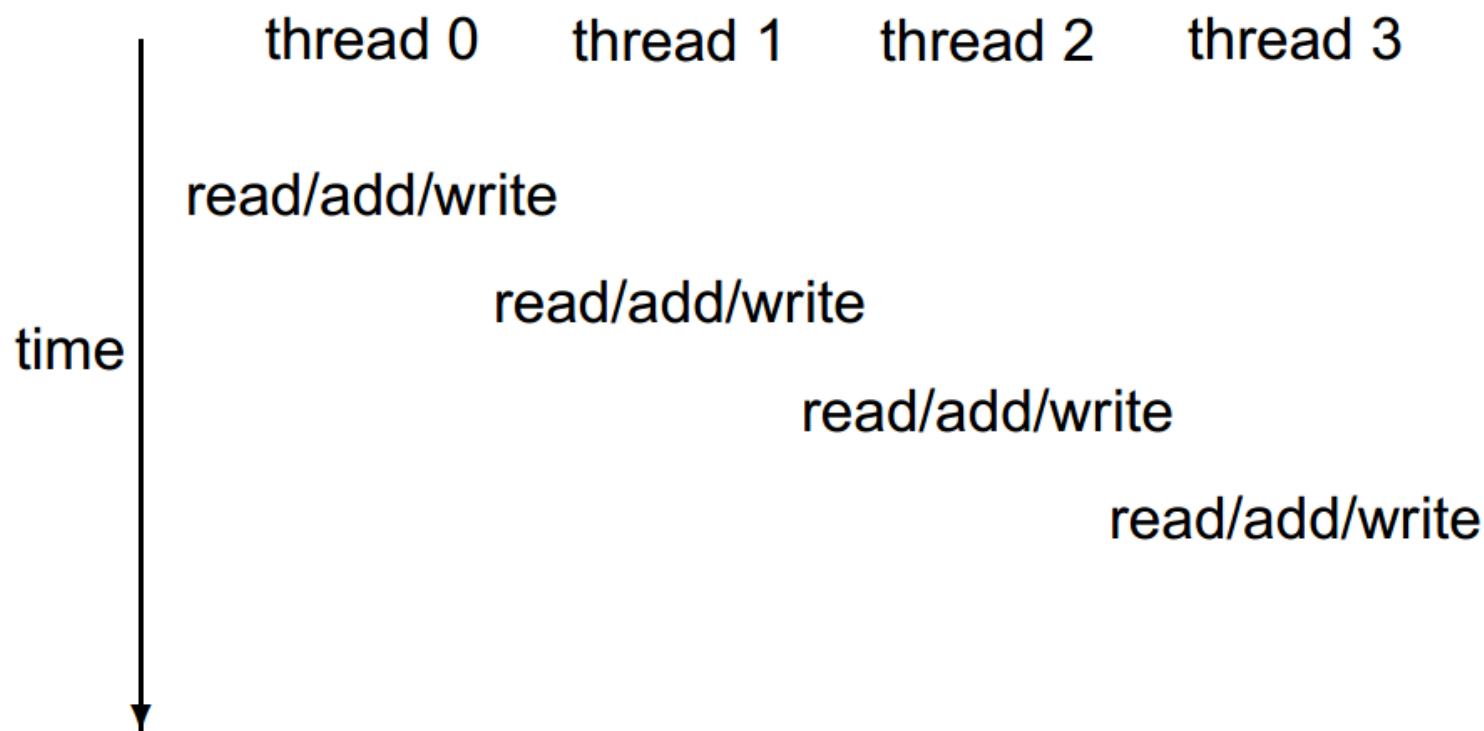
# Atomic operations

- Using standard instructions, multiple threads in the same warp will only update it once.



# Atomic operations

- With atomic instructions, the read/add/write becomes a single operation, and they happen one after the other



# Atomic operations

- Several different atomic operations are supported, almost all only for integers:
  - addition (integers and 32-bit floats)
  - minimum / maximum
  - increment / decrement
  - exchange / compare-and-swap
  
- These are
  - not very fast for data in shared memory (better in next generation Maxwell architecture)
  - only slightly slower for data in device memory (operations performed in L2 cache)

# Atomic operations

## ■ Compare-and-swap:

```
int atomicCAS(int* address, int compare, int val);
```

- if compare **equals** old **value stored at** address **then val is stored instead**
- in either case, routine returns the value of old
- seems a bizarre routine at first sight, but can be very useful for atomic locks
- also can be used to implement 64-bit floating point atomic addition

# Global atomic lock

```
// global variable: 0 unlocked, 1 locked
__device__ int lock=0;

__global__ void kernel(...) {
    ...

    if (threadIdx.x == 0) {
        // set lock
        do {
        } while(atomicCAS(&lock, 0, 1));

        ...

        // free lock
        lock = 0;
    }
}
```

# Global atomic lock

- Problem: when a thread writes data to device memory the order of completion is not guaranteed, so global writes may not have completed by the time the lock is unlocked

```
__global__ void kernel(...) {  
    ...  
  
    if (threadIdx.x == 0) {  
        do {  
        } while(atomicCAS(&lock, 0, 1));  
        ...  
        __threadfence(); // wait for writes to finish  
  
        // free lock  
        lock = 0;  
    }  
}
```

# \_\_threadfence

- `__threadfence_block();`

**wait until all global and shared memory writes are visible to**

- all threads in block

- `__threadfence();`

**wait until all global and shared memory writes are visible to**

- all threads in block

- all threads, for global data

# Atomic addition for double

```
// atomic addition from Jon Cohen at NVIDIA
static double atomicAdd(double *addr, double val)
{
    double old = *addr, assumed;

    do {
        assumed = old;
        old = __longlong_as_double(
            atomicCAS((unsigned long long int*)addr,
                     __double_as_longlong(assumed),
                     __double_as_longlong(val + assumed) )
        );
    } while( assumed != old );

    return old;
}
```



# Summary

- lots of esoteric capabilities – don't worry about most of them
- essential to understand warp divergence – can have a very big impact on performance
- `__syncthreads()` is vital – will see another use of it in next lecture
- the rest can be ignored until you have a critical need – then read the documentation carefully and look for examples in the SDK

# Key reading

## ■ **CUDA Programming Guide, version 5.5:**

- **Section 5.4.2: control flow and predicates**
- **Section 5.4.3: synchronization**
- **Appendix B.5: `__threadfence()` and variants**
- **Appendix B.6: `__syncthreads()` and variants**
- **Appendix B.12: atomic functions**
- **Appendix B.13: warp voting**

# 2D Laplace solver

- Jacobi iteration to solve discrete Laplace equation on a uniform grid:



```
for (int j = 0; j < J; j++) {  
    for (int i = 0; i < I; i++) {  
        id = i + j * I; // 1D memory location  
  
        if (i == 0 || i == I - 1 || j == 0 || j == J - 1)  
            u2[id] = u1[id];  
        else  
            u2[id] = 0.25 * (u1[id - 1] + u1[id + 1]  
                             + u1[id - I] + u1[id + I]);  
    }  
}
```

# 2D Laplace solver


## ■ How do we tackle this with CUDA?

- each thread responsible for one grid point
- each block of threads responsible for a block of the grid
- conceptually very similar to data partitioning in MPI distributed-memory implementations, but much simpler
- (also similar to blocking techniques to squeeze the best cache performance out of CPUs)
- great example of usefulness of 2D blocks and 2D “grid”s


# 2D Laplace solver



# 2D Laplace solver



# 2D Laplace solver



- Each block of threads processes one of these grid blocks, reading in old values and computing new values

# 2D Laplace solver

```
__global__ void lap(int I, int J, const float *__restrict__ u1,
                    float *__restrict__ u2) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int j = threadIdx.y + blockIdx.y * blockDim.y;
    int id = i + j * I;

    if (i == 0 || i == I-1 || j == 0 || j == J-1) {
        u2[id] = u1[id]; // Dirichlet b.c.'s
    }
    else {
        u2[id] = 0.25 * (u1[id - 1] + u1[id + 1]
                        + u1[id - I] + u1[id + I]);
    }
}
```




# 2D Laplace solver

## ■ Assumptions:

- $I$  is a multiple of `blockDim.x`
  - $J$  is a multiple of `blockDim.y`
  - hence grid breaks up perfectly into blocks
- Can remove these assumptions by testing whether  $i, j$  are within grid

# 2D Laplace solver




# 2D Laplace solver

```
__global__ void lap(int I, int J, const float* __restrict__ u1,
                    float* __restrict__ u2) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int j = threadIdx.y + blockIdx.y * blockDim.y;
    int id = i + j * I;

    if (i == 0 || i == I-1 || j == 0 || j == J-1) {
        u2[id] = u1[id]; // Dirichlet b.c.'s
    }
    else if (i < I && j < J) {
        u2[id] = 0.25f * (u1[id - 1] + u1[id + 1]
                        + u1[id - I] + u1[id + I]);
    }
}
```

# 2D Laplace solver

## ■ How does cache function in this application?



- if block size is a multiple of 32 in x -direction, then interior corresponds to set of complete cache lines
- “halo” points above and below are full cache lines too
- “halo” points on side are the problem – each one requires the loading of an entire cache line
- optimal block shape has aspect ratio of roughly 32:1

# 3D Laplace solver

- practical 3
- each thread does an entire line in  $z$ -direction
- $x, y$  dimensions cut up into blocks in the same way as 2D application
- `laplace3d.cu` and `laplace3dkernel.cu` follow same approach described above