

Introduction to CUDA Programming

Lecture 4: libraries and tools

高性能计算机研究中心

CUDA libraries

Originally, NVIDIA planned to provide only one or two maths libraries, but over time these have steadily increased

■ CUBLAS

- basic linear algebra subroutines for dense matrices
- includes matrix-vector and matrix-matrix product
- significant input from Vasily Volkov at UC Berkeley; one routine contributed by Jonathan Hogg from RAL
- with dynamic parallelism on Kepler, it is now possible to call CUBLAS routines from user kernels
- some support for a single routine call to do a “batch” of smaller matrix-matrix multiplications
- also support for using CUDA streams to do a large number of small tasks concurrently

CUDA libraries

- **CUBLAS is a set of routines to be called by user host code:**
 - **helper routines:**
 - memory allocation
 - data copying from CPU to GPU, and vice versa
 - error reporting
 - **compute routines:**
 - matrix-matrix product
 - matrix-vector product
 - **Warning!** Some calls are asynchronous, i.e. the call starts the operation but the host code then continues before it has completed
- **simpleCUBLAS example in SDK is a good example code**

CUDA libraries

■ CUFFT

- Fast Fourier Transform
- 1D, 2D, 3D
- significant input from Satoshi Matsuoka and others at Tokyo Institute of Technology
www.voltaire.com/assets/files/Casestudies/titechcasestudyfinalforSC08.pdf
- has almost all of the variations found in FFTW and other CPU libraries?
- improved performance with Kepler hardware due to new warp shuffle instructions

CUDA libraries

- Like CUBLAS, it is a set of routines called by user host code:
 - helper routines include “plan” construction
 - compute routines perform 1D, 2D, 3D FFTs
 - it supports doing a “batch” of independent transforms, e.g. applying 1D transform to a 3D dataset
 - simpleCUFFT example in SDK

CUDA libraries

■ CUSPARSE

- various routines to work with sparse matrices
- includes sparse matrix-vector and matrix-matrix products
- could be used for iterative solution
- also has solution of sparse triangular system
- note: batched tridiagonal solver is in CUBLAS not CUSPARSE
- contribution from István Reguly (Oxford)

CUDA libraries

■ CURAND

- random number generation
- XORWOW, mrg32k3a, Mersenne Twister and Philox_4x32_10 pseudo-random generators
- Sobol quasi-random generator (with optimal scrambling)
- uniform, Normal, log-Normal, Poisson outputs

■ CUB

- provides a collection of basic building blocks at three levels: device, thread block, warp
- functions include sort, scan, reduction
- Thrust uses CUB for CUDA version of key algorithms

CUDA libraries

■ NPP (NVIDIA Performance Primitives)

- library for imaging and video processing
- includes functions for filtering, JPEG decoding, etc.

■ AmgX (originally named NVAMG)

- library for algebraic multigrid
- see presentation given at Univ. of Warwick:
<http://www2.warwick.ac.uk/fac/sci/dcs/research/pcav/linearsolvers/programme/eatonnvidia.pdf>
or at NVIDIA's 2013 GTC conference:
<http://on-demand.gputechconf.com/gtc/2013/presentations/S3579-High-Performance-Algebraic-Multigrid-Commercial-Apps.pdf>
- available from
<http://developer.nvidia.com/amgx>

CUDA libraries

■ Thrust

- high-level C++ template library with an interface based on the C++ Standard Template Library (STL)
- very different philosophy to other libraries; users write standard C++ code (no CUDA) but get the benefits of GPU parallelisation
- also supports x86 execution
- relies on C++ object-oriented programming; certain objects exist on the GPU, and operations involving them are implicitly performed on the GPU
- I've not used it, but for some applications it can be very powerful – e.g. lots of built-in functions for operations like sort and scan
- also simplifies memory management and data movement

Useful header files

- **dbldbl.h available from**
<https://gist.github.com/seibert/5914108>
Header file for double-double arithmetic for quad-precision (developed by NVIDIA, but published independently under the terms of the BSD license)
- **cuda_complex.h available from**
<https://github.com/jtravs/cudacomplex/blob/master/cudacomplex.hpp>
Header file for complex arithmetic – defines a class and overloaded arithmetic operations. (NVIDIA currently has no plans to adopt this)
- **helper_math.h available in CUDA SDK Defines operator-overloading operations for CUDA intrinsic vector datatypes such as float4**

Other libraries

■ MAGMA

- a new LAPACK for GPUs – higher level numerical linear algebra, layered on top of CUBLAS
- open source – freely available
- Jack Dongarra, Jim Demmel and others

■ FLAME

- similar, but being developed by Robert van de Geijn at UT Austin with various collaborators

Other libraries

■ NAG RNG GPU library

- mrg32k3a and Mersenne Twister pseudo-random generators
- Sobol quasi-random generator
- uniform, Normal, exponential outputs
- Brownian bridge for use with Sobol generator
- www.nag.co.uk/numeric/GPUs/

■ CULAtools

- commercial library, with dense and sparse linear algebra capabilities
- restricted single-precision subset free for academics
- <http://www.culatools.com/>

Other libraries

■ ArrayFire from Acclereyes:

- commercial software, but free for academics and single GPU use
- supports both CUDA and OpenCL execution
- C, C++ and Fortran interfaces
- wide range of functionality including linear algebra, image and signal processing, random number generation, sorting
- www.accelereyes.com/products/arrayfire

■ NVIDIA maintains a webpage with links to a variety of CUDA libraries and other tools:

developer.nvidia.com/cuda-tools-ecosystem

The 7 dwarfs

- Phil Colella, senior researcher at Lawrence Berkeley National Laboratory, talked about “7 dwarfs” of numerical computation in 2004
- expanded to 13 by a group of UC Berkeley professors in a 2006 report: “A View from Berkeley”
www.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.pdf
- key algorithmic kernels in many scientific computing applications
- very helpful to focus attention on HPC challenges and development of libraries and problem-solving environments/frameworks.

The 7 dwarfs

- dense linear algebra
- sparse linear algebra
- spectral methods
- N-body methods
- structured grids
- unstructured grids
- Monte Carlo

Dense linear algebra

- CUBLAS
- MAGMA
- FLAME
- ArrayFire
- CULAtools

Sparse linear algebra

■ iterative solvers:

- some available in Petsc
- others can be implemented using sparse matrix-vector multiplication from CUSPARSE
- NVIDIA has AmgX, an algebraic multigrid library

■ commercial direct solvers:

- Access Analytics (ex-Boeing Computer Services)
- ANSYS/Acceleware

■ non-commercial direct solvers:

- SuperLU project at University of Florida (Tim Davis)
www.cise.ufl.edu/davis/publications_files/qrgpu_paper.pdf
- new project at RAL (Jennifer Scott & Jonathan Hogg)

Spectral methods

- **CUFFT**
 - library provided / maintained by NVIDIA
- **nothing else needed?**

N-body methods

■ OpenMM

- open source package to support molecular modelling, developed at Stanford

■ Fast multipole methods:

- ExaFMM by Yokota and Barba:
<http://www.bu.edu/exafmm/>
- FMM2D by Holm, Engblom, Goude, Holmgren:
<http://user.it.uu.se/stefane/freeware>
- Software by Takahashi, Cecka, Fong, Darve:
onlinelibrary.wiley.com/doi/10.1002/nme.3240/pdf

Structured grids

- lots of people have developed one-off applications
- no great need for a library for single block codes (though possible improvements from “tiling”?)
- multi-block codes could benefit from a general-purpose library, mainly for MPI communication

Unstructured grids

- In addition to GPU implementations of specific codes in CFD community (e.g. Rainald Löhner at GMU / NRL) there are projects to create high-level solutions which others can use for their application codes:
 - Alonso, Darve and others (Stanford)
 - Oxford / Imperial College project has developed OP2, a general-purpose open-source framework based on a previous framework built on MPI
(Case Study 3 on Friday)
- May be other work I'm not aware of

Monte Carlo

- **NVIDIA CURAND library**
- **NAG GPU RNG library**
- **Acclereyes ArrayFire library**
- **some examples in CUDA SDK distribution**
- **nothing else needed except for more output distributions?**

Tools

■ Debugging:

- `cuda-memcheck`

detects array out-of-bounds errors, and mis-aligned device memory accesses – very useful because such errors can be tough to track down otherwise

- `cuda-memcheck --tool racecheck`

this checks for shared memory race conditions:

- **Write-After-Write (WAW):** two threads write data to the same memory location but the order is uncertain
- **Read-After-Write (RAW) and Write-After-Read (WAR):** one thread writes and another reads, but the order is uncertain

Tools

■ Other languages:

- FORTRAN: PGI (Portland Group) CUDA FORTRAN compiler with natural FORTRAN equivalent to CUDA C
- MATLAB: can call kernels directly, or use OOP like Thrust to define MATLAB objects which live on the GPU

<http://www.oerc.ox.ac.uk/projects/cuda-centre-excellence/matlab-gpus>

- Mathematica: similar to MATLAB?

- Python:<http://mathematician.de/software/pycuda>

- R:<http://www.fuzzyl.com/products/gpu-analytics/>
<http://cran.r-project.org/web/views/HighPerformanceComputing.html>

- Haskell:<https://hackage.haskell.org/package/cuda>
<http://hackage.haskell.org/package/accelerate>
<http://chimera.labs.oreilly.com/books/1230000000929/ch06.html>

Tools

■ Other target platforms:

- PGI CUDA FORTRAN and CUDA C compilers are able to generate x86 code for multicore CPUs – this can also produce AVX code for the latest CPU vector units (but I've not tested it)
- CUDA→OpenCL translation for AMD GPUs Swan is a simple tool for porting some programs

<http://gpgpu.org/2010/03/09/swan>

Tools

■ Integrated Development Environments (IDE):

- Nsight Visual Studio edition – NVIDIA plug-in for Microsoft Visual Studio

developer.nvidia.com/nvidia-nsight-visual-studio-edition

- Nsight Eclipse edition – IDE for Linux systems

developer.nvidia.com/nsight-eclipse-edition

- these come with editor, debugger, profiler integration

Tools

■ Visual Profiler:

- standalone NVIDIA software for Linux and Windows systems
- uses hardware counters to collect a lot of useful information
- I think only 1 SMX is instrumented – implicitly assumes the others are behaving similarly
- lots of things can be measured, but a limited number of counters, so it runs the application multiple times if necessary to get full info
- can also obtain instruction counts from command line:
nvprof --metrics "flops_sp, flops_dp" prac2
do nvprof --help for more info on other options

Summary

- active work on all of the dwarfs
- in most cases, significant effort to develop general purpose libraries or frameworks, to enable users to get the benefits without being CUDA experts
- too much going on for one person (e.g. me) to keep track of it all
- NVIDIA maintains a webpage with links to CUDA tools/libraries:
developer.nvidia.com/cuda-tools-ecosystem
- the existence of this eco-system is part of why I think CUDA will remain more used than OpenCL for HPC